

NoSQL and MongoDB

- Objectives
 - Discover the advantages of NoSQL and MongoDB
 - See why the industry is moving towards NoSQL
 - Use PyMongo to access and update MongoDB
 - Learn MongoDB's query language

History: Starting from RDBMSs is an axiom

- Starting from an RDBMS continues to be an axiom of software development.


When was the last time you consciously evaluated alternatives to an RDBMS?

History: Why has SQL persisted so long?

- **Experience:** Industry experience in reliably running RDBMSs is useful.
- **Tooling:** Many many tools speak SQL and understand RDBMSs.
- **ACID:** RDBMSs typically provide app safety via ACID properties.
- **DBAs:** There can be a professional divide between DBAs and developers.
- **Concurrency:** They tame the challenges of concurrency and failure.
- **Integration DBs (1):** RDBMSs have been used in large organizations as integration layers between many enterprise apps.

History: Why has SQL persisted so long?

- The integration database:



History: The industry is moving away from integration dbs

Integration database have issues:


- Their schemas are often more complex (dramatically more so) than databases built for a single application (application databases).
- Complex coordination is required for every little change (innovation through change committees anyone?)
- They drive many important applications, can you scale them?
- They play a central role in “the rise of the DBA”.

History: Application database

- Many large IT groups are moving away from integration databases towards application databases (1) + services and SOA.
 - services (SOA) provide another way to design systems:
 - you access data through the service layer
 - application databases provide data for a single service

History: The industry is moving away from integration dbs

- The application database + services solution:



MongoDB: MongoDB is serious business

- Some companies using MongoDB:
 - <http://www.mongodb.org/about/production-deployments/>

The logo for Craigslist, featuring the word "craigslist" in a blue, lowercase, sans-serif font.The logo for Foursquare, featuring the word "foursquare" in a blue, lowercase, sans-serif font with a trademark symbol.The logo for Developmentor, featuring the word "DEVELOPMENTOR" in a white, uppercase, sans-serif font on a black background, with a white paper airplane icon to the right.The logo for Facebook, featuring the word "facebook" in a blue, lowercase, sans-serif font.The logo for CERN, featuring the word "CERN" in a black, uppercase, sans-serif font on a light gray background.The logo for GitHub, featuring the word "github" in a black, lowercase, sans-serif font, with "SOCIAL CODING" in a smaller, black, uppercase, sans-serif font below it.The logo for SourceForge, featuring the word "sourceforge" in a blue, lowercase, sans-serif font.The logo for bit.ly, featuring the text "bit.ly" in a blue, lowercase, sans-serif font.The logo for Disney, featuring the word "Disney" in a black, stylized, cursive font.The logo for Wordnik, featuring the word "wordnik" in a black, lowercase, sans-serif font, with a red heart symbol replacing the letter "o". Below it is the tagline "All the words." in a small, black, lowercase, sans-serif font.

NoSQL: Document DBs - how do they store data?

```
{
  "_id" : ObjectId("524ca37bd588bf0e4c1ff713"),
  "Name" : "Intensive C++ Training",
  "ActiveCourse" : true,
  "NewCourse" : false,
  "CourseHighlights" : "...",
  "Prerequisites" : "...",
  "Engagements" : [
    {
      "_id" : ObjectId("524ca37bd588bf0e4c1ff714"),
      "CourseId" : ObjectId("524ca37bd588bf0e4c1ff713"),
      "StartDate" : ISODate("2010-03-15T07:00:00Z"),
      "..." : "..."
    },
    {
      "_id" : ObjectId("524ca37bd588bf0e4c1ff715"),
      "CourseId" : ObjectId("524ca37bd588bf0e4c1ff713"),
      "StartDate" : ISODate("2011-04-11T07:00:00Z"),
      "..." : "..."
    }
  ],
  "CourseAliases" : [],
  "UrlPath" : "intensive-c++-training"
}
```

Scaling: RDBMS

- Typically vertical scaling



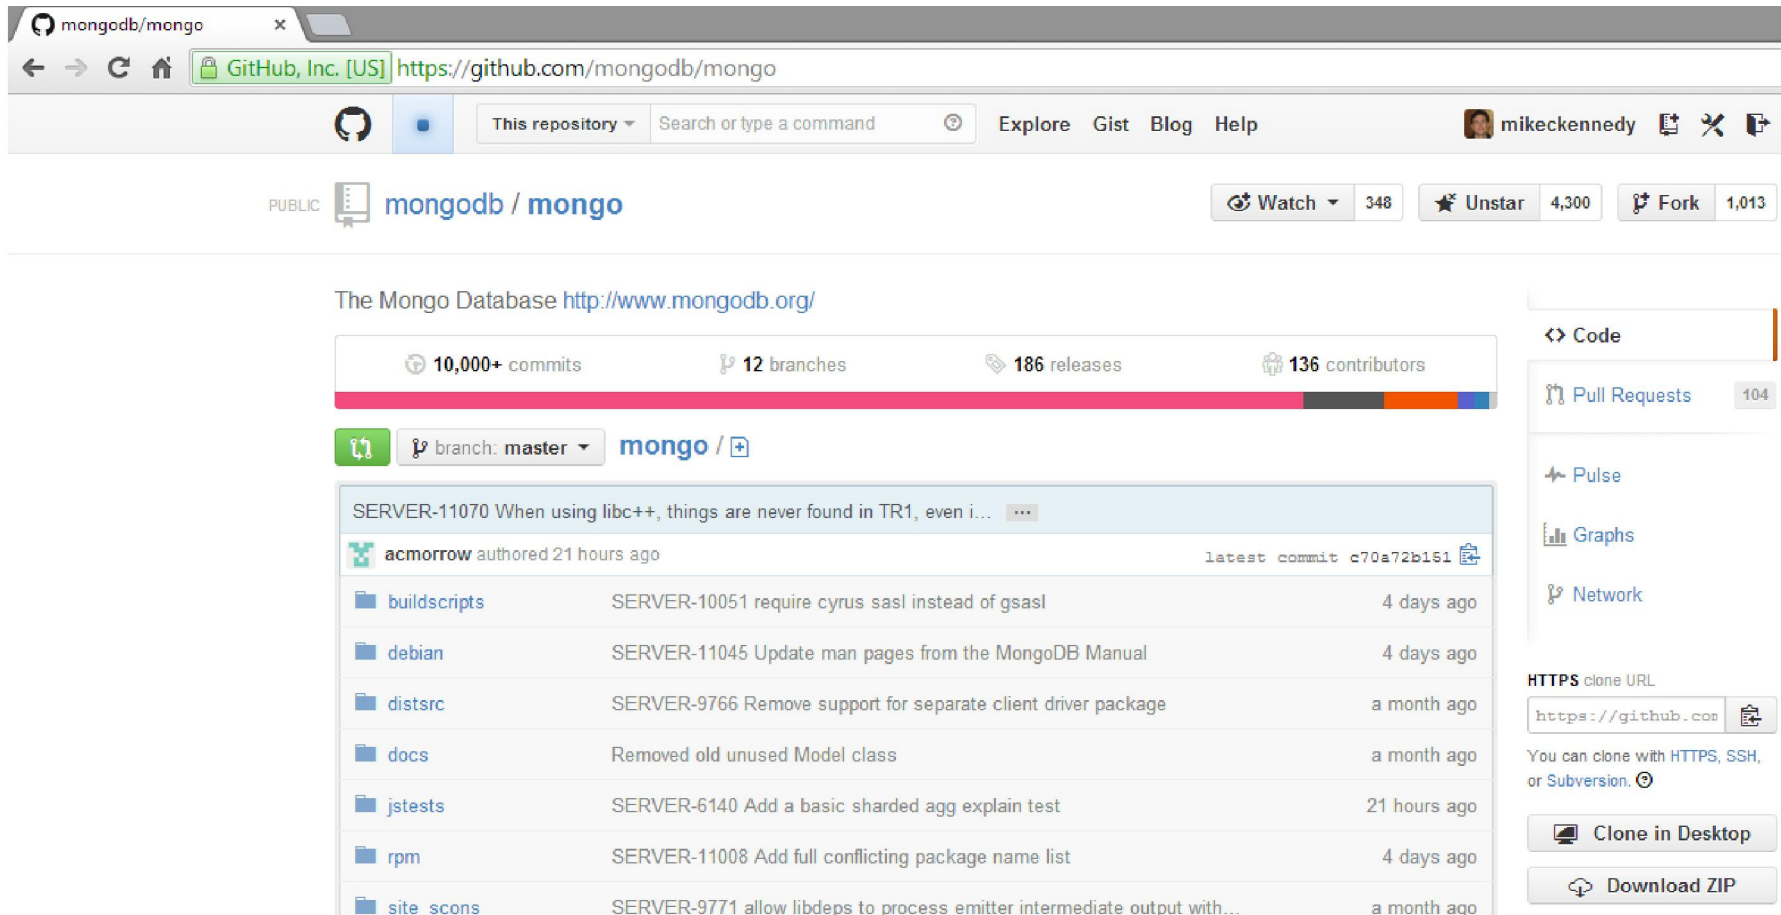
Start here
\$ / perf.



Scale here
\$\$\$\$ / perf.

MongoDB: Open source

- Source code (server and drivers) on github
- Free download (pay for support)



The screenshot shows the GitHub repository for MongoDB. The browser address bar displays the URL `https://github.com/mongodb/mongo`. The repository is public and has 348 watchers, 4,300 stars, and 1,013 forks. The repository name is `mongodb / mongo`. The main content area shows the commit history for the `master` branch. The latest commit is by `acmorrow` 21 hours ago, with the commit hash `c70a72b151`. The commit message is `SERVER-11070 When using libc++, things are never found in TR1, even i...`. The commit history table lists several other commits with their messages and timestamps.

Commit	Message	Time
acmorrow	SERVER-11070 When using libc++, things are never found in TR1, even i...	21 hours ago
buildscripts	SERVER-10051 require cyrus sasl instead of gsasl	4 days ago
debian	SERVER-11045 Update man pages from the MongoDB Manual	4 days ago
distsrc	SERVER-9766 Remove support for separate client driver package	a month ago
docs	Removed old unused Model class	a month ago
jstests	SERVER-6140 Add a basic sharded agg explain test	21 hours ago
rpm	SERVER-11008 Add full conflicting package name list	4 days ago
site_scons	SERVER-9771 allow libdeps to process emitter intermediate output with...	a month ago

On the right side, there are links for `Code`, `Pull Requests` (104), `Pulse`, `Graphs`, and `Network`. At the bottom, there are buttons for `Clone in Desktop` and `Download ZIP`.

MongoDB: Getting MongoDB

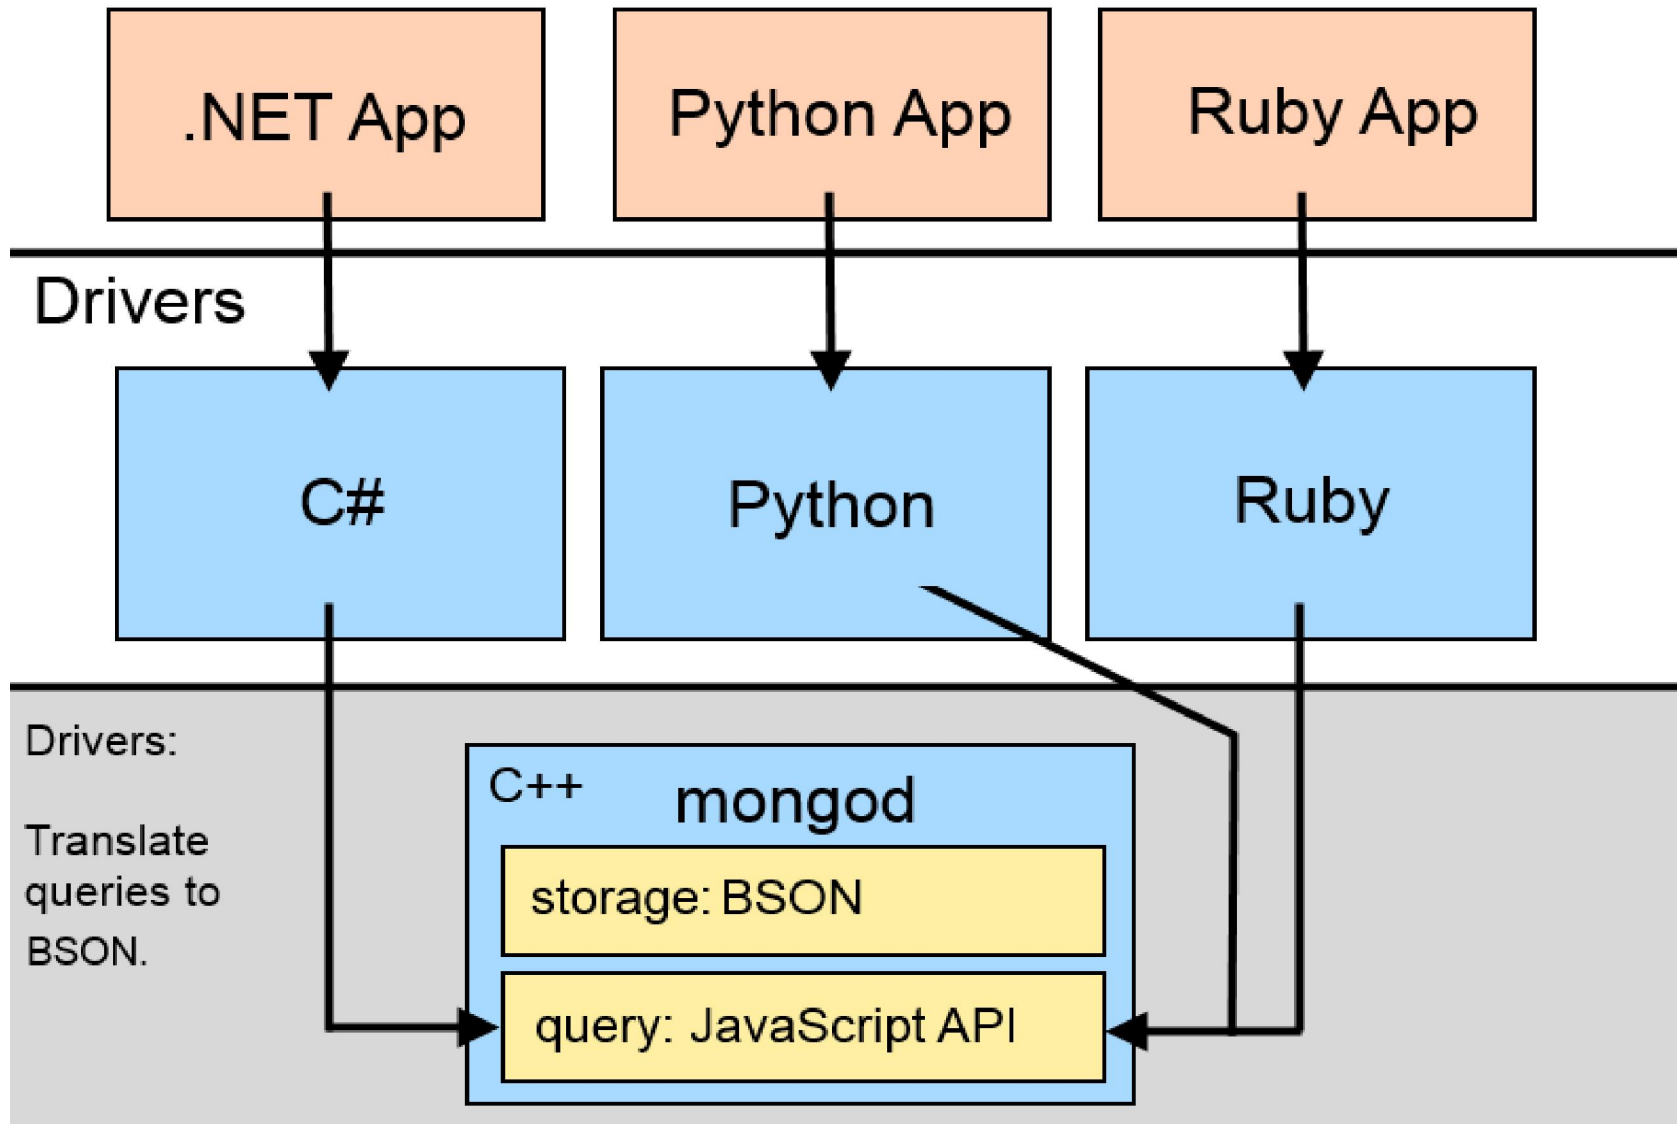
- Download your OS's version here
 - free, no registration required,
 - <http://www.mongodb.org/downloads> (64-bit is preferred)
- Install MongoDB by decompressing it
 - you may want to set it to start up as a Windows Server or system daemon
- Create a data folder and log folder
- Start the server

Estimated time: First time, less than 10 min, with experience, less than 3 min.

MongoDB: Management tools

- **Robomongo**
 - Free / Open-source, Windows / OS X / Linux
 - <http://robomongo.org/>
- **Mongo.exe** – MongoDB's native shell
 - Free / Open-source, Windows / OS X / Linux

The internals of MongoDB



PyMongo

- MongoDB has an official Python driver
 - **pymongo**: <https://pypi.python.org/pypi/pymongo>
- Tutorials and documentation from MongoDB
 - Python Language Center
<http://docs.mongodb.org/ecosystem/drivers/python/>
- Open-source on Github
 - <https://github.com/mongodb/mongo-python-driver>
- Supports
 - Python 3 and Python 2
 - Windows, OS X, Linux
- Installing pymongo from the installers on PyPI is preferred

PyMongo [connecting]

import pymongo
(after installing)

Start with `pymongo.MongoClient()`

```
import pymongo  
  
mongo = pymongo.MongoClient()  
  
db = mongo.DM_Books  
collection = db.Publisher
```

MongoClient uses dynamic programming to access DBs and collections

Note: If **DM_Books** or **Publisher** does not exist, this is how you create them.

Querying for documents: find, find_one

- For simple queries, we use find and specify prototype dictionaries.

```
set = db.Collection.find( { '_id': 1 } )
# Think: SELECT * FROM Collection WHERE _id = 1

set = db.Collection.find( {'username': "mkennedy" } )
# Think: SELECT * FROM Collection WHERE username = 'mkennedy'

doc = db.Collection.find_one(
    {'username': "pierre", 'passwordhash': "B8E28A21" } )
# Think: SELECT TOP(1) * FROM Collection
#         WHERE username = 'pierre' AND ' passwordhash = 'B8E28A21'
```

Querying for documents: results

The return value from find is a cursor.



```
query = db.Collection.find( {'Category': 'NoSQL'} )

# Get the number of records with count
numOfRecords = query.count()

# Pull all documents into memory as a list of dicts
memoryList = list(query)

# Stream documents to the app via for/in loops
for doc in query:
    print(doc['Title'])
```

PyMongo [viewing results]

- Highly-nested data can be pretty printed for readability

```
import pprint

doc = db.Collection.find_one( {'Category': 'NoSQL'} );

pprint.pprint(doc)
# prints

{
  'Author': 'Joe Vitale',
  'ISBN': '0759614318',
  'Published': datetime.datetime(2001, 1, 1, 8, 0),
  'Publisher': ObjectId('5258672c3a93bb21980ffa8d'),
  'Category': 'NoSQL',
  'Title': 'Spiritual Databases: A ...',
  '_id': ObjectId('525867633a93bb2198137c81')
}
```

Operators: Introduction

- How would you express this as a prototypical dictionary?
 - `SELECT * FROM Users WHERE RegistrationDate > @date`
- You cannot, which is why we need **\$operators**. e.g.

```
db.Users.find( { 'registrationDate': { '$gt': date } } )
```

Operators: Inequalities and existence

```
db.Users.find( { 'registrationDate': { '$gt': date } } )
```

- **\$gt** - greater than
- **\$gte** - greater than or equal to
- **\$lt** - less than
- **\$lte** - less than or equal to
- **\$ne** - not equal
- **\$exists** - the field exists in this document

Operators: Inequalities and existence

- Operator example, find non-null email addresses:

```
db.Users.find( { 'email': { '$ne': None } } )
```

Operators: \$or and \$and

- Often you need to combine two filters using OR or AND:

```
db.Users.find( { '$and':  
  [  
    { 'email': { '$ne': None } },  
    { 'registrationDate': { '$gt': date } }  
  ]  
} )
```

- Note: \$and and \$or expect an array of conditions.

Operators: inside arrays (\$in, \$all)

- It is very common to traverse a weak foreign key using two queries.

```
{ // category
  _id: "science",
  bookIds: [1, 5, 93, 20, 11]
  // more items
}
```

```
cat = db.Categories.find( { '_id': 'science' } )
books = db.Books.find( { '_id': { '$in': cat.bookIds } } )
```


Updating documents: Entire documents

- We can treat MongoDB as an ORM.
 - get document
 - make changes
 - save document back to DB

```
user = db.Users.find_one( { '_id': 72 } )  
user['hasPaid'] = True  
user['expirationDate'] = newEndDate  
db.Users.save(user)
```

Updating documents: By field, atomically

- **\$set** atomically updates the document without retrieving it.

```
db.Users.update(  
  { '_id': 72 },  
  { '$set':  
    {  
      'hasPaid': True,  
      'expirationDate' = newEndDate  
    }  
  } )
```

Updating multiple records

- By default, the `update()` method updates a single document. If the `multi` option is set to `true`, the method updates all documents that match the query criteria.

```
db.Users.update(  
  { 'hasPaid': True },  
  { '$set': { expirationDate = newEndDate } },  
  { 'multi': True }  
)
```

Warning: This goes against your intuition.
SQL updates everything that matches by default.

Deleting / removing documents

- Delete documents with `db.collection.remove`:

```
# remove all non paying users.  
db.Users.remove( { 'hasPaid': False} )
```

Note: For large deletion operations, it may be more efficient to copy the documents that you want to keep to a new collection and then use `drop()` on the original collection.

Summary

- MongoDB is a cluster-friendly, scalable database
- Simpler programming models and application DBs are leading developers down the NoSQL path
- PyMongo is the official MongoDB driver for Python
- MongoDB has a prototypical document-based query language